

sahib-holistic-credit-card-fraud-detection_v2

April 18, 2024

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OrdinalEncoder

from typing import List, Tuple
from collections import namedtuple
from sklearn.base import BaseEstimator, TransformerMixin #, clone
from sklearn.metrics import accuracy_score, classification_report, roc_auc_score

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB, CategoricalNB
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from hyperopt import hp, STATUS_OK, tpe, fmin, Trials
import warnings
import regex as re

import joblib
import os
import json

from sklearn.feature_selection import RFE

from sklearn.model_selection import cross_val_score

pd.set_option('display.max_columns', None)
warnings.filterwarnings('ignore')

[1]: # setting up directories for better organizing trained models, reports and
↳metrics
MODEL_DIR = './models/'
```

```
DATASET_DIR = './datasets/'
REPORTS_DIR = './metrics/performance'
```

1 Loading Base Dataset

```
[3]: # load the raw dataset
base_dataset = pd.read_csv('/kaggle/input/fraud-detection/fraudTrain.csv',
    ↳na_values= ['?', '$', '#', 'unknown', 'Unknown']).iloc[:, 1:]
print(f'Columns : {base_dataset.columns.tolist()}')
print(f'Shape : {base_dataset.shape}')
```

```
Columns : ['trans_date_trans_time', 'cc_num', 'merchant', 'category', 'amt',
'first', 'last', 'gender', 'street', 'city', 'state', 'zip', 'lat', 'long',
'city_pop', 'job', 'dob', 'trans_num', 'unix_time', 'merch_lat', 'merch_long',
'is_fraud']
Shape : (1296675, 22)
```

```
[4]: # look at the label distribution - inspecting for label imbalance
base_dataset.is_fraud.value_counts()
```

```
[4]: is_fraud
0      1289169
1         7506
Name: count, dtype: int64
```

```
[5]: # balancing the base dataset with equal number of fraud and non-fraud datapoints
not_fraud = base_dataset[base_dataset['is_fraud'] == 0].sample(n = 1500,
    ↳random_state = 1)
is_fraud = base_dataset[base_dataset['is_fraud'] == 1].sample(n = 1500,
    ↳random_state = 1)
del base_dataset
base_dataset_balanced = pd.concat([not_fraud, is_fraud], axis = 0).sample(n =
    ↳3000).reset_index(drop = True)

# TODO : REMOVE COMMENT BELOW
# base_dataset_balanced = pd.read_csv('./datasets/balanced_base_dataset.csv')
# base_dataset_balanced.head(2)

# segregating transaction (t_cols) and demographic-related (d_cols) columns
t_cols =
    ↳'trans_date_trans_time,cc_num,merchant,category,amt,trans_num,unix_time,merch_lat,merch_long'
    ↳split(',')
d_cols =
    ↳'first,last,gender,street,city,state,zip,lat,long,city_pop,job,dob,is_fraud'
    ↳split(',')
```

2 Loading Enrichment Dataset

```
[6]: # both base and enrichment dataset contains 3000 rows
nrows = 3000

#demog1 - adult census
e1 = pd.read_csv('/kaggle/input/adult-census-income/adult.csv', nrows = nrows ,
    ↪na_values= ['?', '$', '#', 'unknown', 'Unknown'])

# demog2 - us consumer finance complaints
e2 = pd.read_csv('/kaggle/input/us-consumer-finance-complaints/
    ↪consumer_complaints.csv', nrows = nrows, na_values= ['?', '$', '#',
    ↪'unknown', 'Unknown'])

# # demog3 - amazon consumer behaviour dataset
# e3 = pd.read_csv('/kaggle/input/amazon-consumer-behaviour-dataset/Amazon
    ↪Customer Behavior Survey.csv', nrows= nrows, na_values= ['?', '$', '#',
    ↪'unknown', 'Unknown'])

# demog3 - ecommerce-behavior-data-from-multi-category-store
e3 = pd.read_csv('/kaggle/input/
    ↪ecommerce-behavior-data-from-multi-category-store/2019-Nov.csv', nrows=
    ↪nrows, na_values= ['?', '$', '#', 'unknown', 'Unknown'])

# trans1 - online payments fraud detection
e4 = pd.read_csv('/kaggle/input/online-payments-fraud-detection-dataset/
    ↪PS_20174392719_1491204439457_log.csv', nrows = nrows, na_values= ['?', '$',
    ↪'#', 'unknown', 'Unknown'])

# trans2 - credit card transactions
e5 = pd.read_csv('/kaggle/input/credit-card-transactions/
    ↪credit_card_transactions-ibm_v2.csv', nrows = nrows, na_values= ['?', '$',
    ↪'#', 'unknown', 'Unknown'])

# trans3 - customer shopping dataset
e6 = pd.read_csv('/kaggle/input/customer-shopping-dataset/
    ↪customer_shopping_data.csv', nrows = nrows, na_values= ['?', '$', '#',
    ↪'unknown', 'Unknown'])

e_df = pd.concat([e1, e2, e3, e4, e5, e6], axis =1) # e1.join([e2, e3, e4, e5,
    ↪e6])
# e_df = pd.read_csv('./datasets/enrichment_data.csv')
print(f'Shape of the raw enrichment dataset : {e_df.shape}')
```

Shape of the raw enrichment dataset : (3000, 78)

```
[8]: e_df.head(2)
```

```
[8]:  workclass_DE education_DE marital.status_DE relationship_DE race_DE \
0      NaN      HS-grad      Widowed      Not-in-family      White
1      Private      HS-grad      Widowed      Not-in-family      White

      hours.per.week_DE income_DE      company_DE tags_TE product_id_TE \
0      40      <=50K      U.S. Bancorp      NaN      1003461
1      18      <=50K      Wells Fargo & Company      NaN      5000088

      category_id_TE      category_code_TE brand_TE user_id_TE \
0  2053013555631882655      electronics.smartphone      xiaomi      520088904
1  2053013566100866035      appliances.sewing_machine      janome      530496790

      user_session_TE type_TE      Use Chip_TE MCC_TE \
0  4d3b30da-a5e4-49df-b1a8-ba5943f1dd33      PAYMENT      Swipe Transaction      5300
1  8e5f4f83-366c-4f70-860e-ca7417414283      PAYMENT      Swipe Transaction      5411

      shopping_mall_TE
0      Kanyon
1      Forum Istanbul
```

2.0.1 Cleaning the enrichment dataset

```
[15]: # remove any traces of non-transaction related data
# segregate features into demographic/transactional
# eliminate any repeating features in enriched d_col or t_cols
cols_to_remove = ['amount', 'price', 'Amount', 'Zip', 'unix_time', 'Merchant',
↳City',
↳'Merchant State', 'Year', 'Month', 'Day', 'Time',
↳'invoice_date',
↳'isFraud', 'isFlaggedFraud', 'Is Fraud?', 'invoice_no',
↳'nameOrig',
↳'step', 'oldbalanceOrig', 'newbalanceOrig', 'nameDest',
↳'oldbalanceDest',
↳'newbalanceDest', 'User', 'Card', 'Errors?', 'customer_id',
↳'payment_method',
↳'product', 'sub_product', 'education.num', 'capital.gain',
↳'captial.loss',
↳'date_received', 'issue', 'sub_issue',
↳'consumer_complaint_narrative',
↳'company_public_response', 'state', 'zipcode',
↳'consumer_consent_provided',
↳'submitted_via', 'date_sent_to_company',
↳'company_response_to_consumer',
↳'timely_response', 'consumer_disputed?', 'complaint_id',
↳'event_time', 'event_type',
]
```

```

# filtering 1 - enriched dataset alignment to transaction -related data
e_df = e_df[e_df.columns[~e_df.columns.isin(cols_to_remove)]]

# filtering 2 - segregating refined enriched columns to d_cols and t_cols
d_cols_enriched = ['workclass', 'fnlwgt', 'education', 'marital.status',
                   'occupation', 'relationship', 'race', 'sex', 'hours.per.
↪ week',
                   'native.country', 'income', 'company', 'gender', 'age']

t_cols_enriched = ['tags', 'product_id', 'category_id', 'category_code',
                   'brand', 'user_id', 'user_session', 'type', 'Use Chip',
                   'Merchant Name', 'MCC', 'category', 'quantity',
↪ 'shopping_mall']

# filtering 3 - toning down any overlapping column present in base d_cols and
↪ t_cols from enriched d_cols and t_cols
d_cols_to_remove = ['gender', 'native.country', 'fnlwgt', 'occupation', 'age',
↪ 'sex']
t_cols_to_remove = ['Merchant Name', 'category', 'quantity']

# final enrichment cols
d_cols_enriched = [col for col in d_cols_enriched if col not in
↪ d_cols_to_remove]
t_cols_enriched = [col for col in t_cols_enriched if col not in
↪ t_cols_to_remove]

# final enriched df contained enriched d_cols and t_cols
e_df = e_df[d_cols_enriched + t_cols_enriched]

tagged_t_cols = [f'{col}_TE' for col in t_cols_enriched]
tagged_d_cols = [f'{col}_DE' for col in d_cols_enriched]

# # final enriched dataset with columns tagged depending on demographic or
↪ transactional
e_df.columns = tagged_d_cols + tagged_t_cols
print(f'Shape of the cleaned enrichment dataset : {e_df.shape}')

```

Shape of the cleaned enrichment dataset : (3000, 19)

```
[12]: e_df.head(2)
```

```

[12]:  workclass_DE education_DE marital.status_DE relationship_DE race_DE \
0      NaN      HS-grad      Widowed      Not-in-family      White
1      Private      HS-grad      Widowed      Not-in-family      White

      hours.per.week_DE income_DE      company_DE tags_TE product_id_TE \

```

0	40	<=50K	U.S. Bancorp	NaN	1003461
1	18	<=50K	Wells Fargo & Company	NaN	5000088

	category_id_TE	category_code_TE	brand_TE	user_id_TE	\
0	2053013555631882655	electronics.smartphone	xiaomi	520088904	
1	2053013566100866035	appliances.sewing_machine	janome	530496790	

	user_session_TE	type_TE	Use Chip_TE	MCC_TE	\
0	4d3b30da-a5e4-49df-b1a8-ba5943f1dd33	PAYMENT	Swipe Transaction	5300	
1	8e5f4f83-366c-4f70-860e-ca7417414283	PAYMENT	Swipe Transaction	5411	

	shopping_mall_TE
0	Kanyon
1	Forum Istanbul

3 Enrichment

Joining the base dataset and cleaned enrichment dataset to obtain the dataset that this analysis will be based on.

```
[16]: #creating master dataset by joining the base dataset and enrichment dataset
master_dataset = base_dataset_balanced[d_cols[:-1] + t_cols].join(e_df)
master_dataset.head(2)
```

```
[16]:
```

	first	last	gender	street	city	state	zip	\
0	Kimberly	Rice	F	63991 Destiny Rue Apt. 651	Tyler	TX	75703	
1	Ellen	Carrillo	F	9431 Amanda Mills	Odessa	MO	64076	

	lat	long	city_pop	job	dob	\
0	32.2768	-95.3031	144160	Sports development officer	1984-05-04	
1	38.9829	-93.9757	9512	Clinical research associate	1972-12-31	

	trans_date_trans_time	cc_num	merchant	\
0	2019-06-19 01:18:47	6506116513503136	fraud_Goodwin-Nitzsche	
1	2019-10-14 19:31:20	676314217768	fraud_Kihn, Abernathy and Douglas	

	category	amt	trans_num	unix_time	\
0	grocery_pos	347.88	7266fcbb0c6dedcff4aaca922fb3aa66	1340068727	
1	shopping_net	15.38	e1d3adfb522e1f1476f0b71a022be2ce	1350243080	

	merch_lat	merch_long	is_fraud	workclass_DE	education_DE	\
0	32.063337	-94.562374	1	NaN	HS-grad	
1	39.142095	-93.700393	0	Private	HS-grad	

	marital.status_DE	relationship_DE	race_DE	hours.per.week_DE	income_DE	\
0	Widowed	Not-in-family	White	40	<=50K	
1	Widowed	Not-in-family	White	18	<=50K	

	company_DE	tags_TE	product_id_TE	category_id_TE	\
0	U.S. Bancorp	NaN	1003461	2053013555631882655	
1	Wells Fargo & Company	NaN	5000088	2053013566100866035	

	category_code_TE	brand_TE	user_id_TE	\
0	electronics.smartphone	xiaomi	520088904	
1	appliances.sewing_machine	janome	530496790	

	user_session_TE	type_TE	Use Chip_TE	MCC_TE	\
0	4d3b30da-a5e4-49df-b1a8-ba5943f1dd33	PAYMENT	Swipe Transaction	5300	
1	8e5f4f83-366c-4f70-860e-ca7417414283	PAYMENT	Swipe Transaction	5411	

	shopping_mall_TE
0	Kanyon
1	Forum Istanbul

```
[17]: print(f'Shape of the master dataset : {master_dataset.shape}')
```

Shape of the master dataset : (3000, 41)

4 Preprocessing the master dataset

```
[17]: # check for missing values
missing_vals = round(master_dataset.isna().sum()/master_dataset.shape[0] * 100,
↳1)
missing_cols = missing_vals[missing_vals > 0].index.tolist()
print('Columns with missing values in %')
missing_vals[missing_vals > 0]
```

Columns with missing values in %

```
[17]: workclass_DE      3.4
tags_TE      85.9
category_code_TE  36.9
brand_TE      18.6
dtype: float64
```

4.1 Segregate columns into categorical and numerical

This is done to devise a strategy for handling missing values

```
[18]: # categorize columns into numerical and categorical cols

cat_cols = ['first', 'last', 'gender', 'street', 'city',
            'state', 'zip', 'job', 'cc_num', 'merchant',
            'category', 'trans_num', 'workclass_DE', 'education_DE',
```

```

        'marital.status_DE', 'relationship_DE', 'race_DE'
        , 'hours.per.week_DE', 'income_DE', 'company_DE',
        'tags_TE', 'product_id_TE', 'category_id_TE', 'category_code_TE',
        'brand_TE', 'user_id_TE', 'user_session_TE', 'type_TE',
        'Use Chip_TE', 'MCC_TE', 'shopping_mall_TE']
num_cols = [col for col in master_dataset.columns.tolist() if (col not in
    ↪ cat_cols) and (col != 'is_fraud')]

# are all missing cols categorical ?
print([True if col in cat_cols else False for col in missing_cols])

```

[True, True, True, True]

4.2 Handle missing values

Strategy chosen is as follows :

1. If percentage of missing values is :
 - a. < 40 %, then drop the column
 - b. = 40 %, then impute the column using most frequent value

```

[19]: # DROP < 40 %

# impute cols with missing values less than 40 % and dropping those with
    ↪ greater than that
master_dataset = base_dataset_balanced[d_cols[:-1] + t_cols].join(e_df)
master_dataset.drop(columns = ['tags_TE'], inplace = True)
# remove dropped column from cat_cols
cat_cols.remove('tags_TE')
tagged_t_cols.remove('tags_TE')

# IMPUTE >= 40 %
#define sequence of ops on cat cols -> complex or composite transformer
cat_col_transformer = Pipeline([
    ('mean_imputer', SimpleImputer(strategy='most_frequent')),
    # ('ordinal_encoder', OrdinalEncoder(dtype=int))
],)

cat_col_preprocessor = ColumnTransformer([
    ('cat_transformer', cat_col_transformer, cat_cols)
],
remainder='drop')

cat_col_encoding_transformer = ColumnTransformer([
    ('cat_encoder', OrdinalEncoder(dtype = int), cat_cols)
],
remainder = 'drop')

```



```

prepped_cat_df = pd.DataFrame(cat_col_preprocessor.
    ↪fit_transform(master_dataset), columns = cat_cols)
prepped_cat_df.head(2)

```

```

[19]:      first      last gender      street      city state      zip \
0  Kimberly      Rice      F 63991 Destiny Rue Apt. 651      Tyler      TX 75703
1    Ellen Carrillo      F      9431 Amanda Mills      Odessa      MO 64076

      job      cc_num \
0  Sports development officer 6506116513503136
1  Clinical research associate      676314217768

      merchant      category \
0      fraud_Goodwin-Nitzsche      grocery_pos
1  fraud_Kihn, Abernathy and Douglas      shopping_net

      trans_num workclass_DE education_DE \
0  7266fcbb0c6dedcff4aaca922fb3aa66      Private      HS-grad
1  e1d3adfb522e1f1476f0b71a022be2ce      Private      HS-grad

      marital.status_DE relationship_DE race_DE hours.per.week_DE income_DE \
0      Widowed      Not-in-family      White      40      <=50K
1      Widowed      Not-in-family      White      18      <=50K

      company_DE product_id_TE      category_id_TE \
0      U.S. Bancorp      1003461      2053013555631882655
1  Wells Fargo & Company      5000088      2053013566100866035

      category_code_TE brand_TE user_id_TE \
0      electronics.smartphone      xiaomi      520088904
1  appliances.sewing_machine      janome      530496790

      user_session_TE      type_TE      Use Chip_TE MCC_TE \
0  4d3b30da-a5e4-49df-b1a8-ba5943f1dd33      PAYMENT      Swipe Transaction      5300
1  8e5f4f83-366c-4f70-860e-ca7417414283      PAYMENT      Swipe Transaction      5411

      shopping_mall_TE
0      Canyon
1  Forum Istanbul

```

4.3 Exploratory Data Analysis

This is done to prepare the dataset columns like analyzing which columns needs to be kept and removed, which needs to be transformed for better usability and visualization.

```
[16]: #combine numerical columns and preprocessed categorical df
eda_df = prepped_cat_df.join(master_dataset[num_cols + ['is_fraud']])
eda_df.head(2)
```

```
[16]:      first      last gender      street      city state      zip \
0  Kimberly      Rice      F 63991 Destiny Rue Apt. 651  Tyler      TX 75703
1    Ellen Carrillo      F      9431 Amanda Mills  Odessa      MO 64076

      job      cc_num \
0  Sports development officer  6506116513503136
1  Clinical research associate      676314217768

      merchant ...      lat      long city_pop \
0      fraud_Goodwin-Nitzsche ... 32.2768 -95.3031 144160
1  fraud_Kihn, Abernathy and Douglas ... 38.9829 -93.9757 9512

      dob trans_date trans_time      amt      unix_time merch_lat merch_long \
0 1984-05-04 2019-06-19 01:18:47 347.88 1340068727 32.063337 -94.562374
1 1972-12-31 2019-10-14 19:31:20 15.38 1350243080 39.142095 -93.700393

      is_fraud
0          1
1          0

[2 rows x 40 columns]
```

```
[17]: # list of unwanted cols as they are unrelated to the analysis
unwanted_cols = ['first', 'last', 'street', 'unix_time', 'category_id_TE',
↪ 'category_code_TE'] # city - can be included

# columns that needs to change form for use in analysis
form_change_cols = ['dob', 'trans_date_trans_time']

# checking for columns with cardinality = 1 - to remove if any
for col in eda_df.columns[:-1]:
    if eda_df[col].unique().size <= 1 :
        print(f'Unique values in {col} : {eda_df[col].unique().size}')
```

```
[18]: # remove redundant cols
eda_df = eda_df[eda_df.columns[~eda_df.columns.isin(unwanted_cols)]]
print(f'Shape of the dataset after removing redundant columns : {eda_df.shape}')
```

Shape of the dataset after removing redundant columns : (3000, 34)

```
[19]: eda_df.head(2)
```

```
[19]:  gender      city state      zip      job      cc_num \
0      F      Tyler      TX  75703  Sports development officer  6506116513503136
1      F      Odessa      MO  64076  Clinical research associate      676314217768

      merchant      category \
0      fraud_Goodwin-Nitzsche  grocery_pos
1  fraud_Kihn, Abernathy and Douglas  shopping_net

      trans_num workclass_DE ... shopping_mall_TE \
0  7266fcbb0c6dedcff4aaca922fb3aa66      Private ...      Canyon
1  e1d3adfb522e1f1476f0b71a022be2ce      Private ...      Forum Istanbul

      lat      long city_pop      dob trans_date_trans_time      amt \
0  32.2768 -95.3031  144160  1984-05-04  2019-06-19 01:18:47  347.88
1  38.9829 -93.9757  9512  1972-12-31  2019-10-14 19:31:20  15.38

      merch_lat merch_long is_fraud
0  32.063337 -94.562374  1
1  39.142095 -93.700393  0

[2 rows x 34 columns]
```

```
[20]: eda_df['merchant'].apply(lambda x : x.split('_')[-1])
```

```
[20]: 0      Goodwin-Nitzsche
1      Kihn, Abernathy and Douglas
2      Auer-Mosciski
3      Botsford Ltd
4      Barrows PLC

...
2995      Bechtelar-Rippin
2996      Lockman, West and Runte
2997      Yost, Block and Koepp
2998      Kihn-Schuster
2999      Willms, Kris and Bergnaum
Name: merchant, Length: 3000, dtype: object
```

```
[21]: # handling columns that changes forms
eda_df['age'] = pd.Timestamp.now().year - pd.
    ↳to_datetime(eda_df[form_change_cols[0]]).dt.year
eda_df['trans_datetime'] = pd.to_datetime(eda_df[form_change_cols[1]])
eda_df['trans_year'] = eda_df['trans_datetime'].dt.year
eda_df['trans_month'] = eda_df['trans_datetime'].dt.month_name()
eda_df['merchant'] = eda_df['merchant'].apply(lambda x : x.split('_')[-1])
eda_df.drop(columns = form_change_cols + ['trans_datetime'], inplace = True)
eda_df.head(2)
```

```
[21]: gender    city state    zip                job                cc_num \
0      F    Tyler    TX    75703    Sports development officer    6506116513503136
1      F    Odessa    MO    64076    Clinical research associate        676314217768

            merchant    category \
0            Goodwin-Nitzsche    grocery_pos
1    Kihn, Abernathy and Douglas    shopping_net

            trans_num workclass_DE ...    lat    long \
0    7266fcbb0c6dedcff4aaca922fb3aa66    Private ...    32.2768 -95.3031
1    e1d3adfb522e1f1476f0b71a022be2ce    Private ...    38.9829 -93.9757

    city_pop    amt    merch_lat merch_long is_fraud age trans_year trans_month
0    144160    347.88    32.063337 -94.562374    1    40    2019    June
1     9512    15.38    39.142095 -93.700393    0    52    2019    October

[2 rows x 35 columns]
```

```
[22]: # processing job description text to include designations alone
suffixes = ['opy', 'ake', 'ub', 'and', 'iate', 'ary',
            'son', 'sta', 'ath', 'geon', 'ner', 'wer',
            'fer', 'mer', 'nal', 'rew', 'ief', 'ler',
            'ker', 'der', 'cer', 'yer',
            'ter', 'ist', 'per', 'ger',
            'gner', 'tive', 'ect', 'eer', 'ian',
            'ant', 'her', 'ot', 'or', 'ser',
            'dic', 'rse', 'ney', 'yst', 'per']
re_pattern_v2 = '|'.join([r"\b(\w+%s\b)" % suffix for suffix in suffixes])
re_pattern_v2

eda_df['job'] = eda_df['job'].str.extract(re_pattern_v2, flags=re.IGNORECASE,
    ↪ expand = False).fillna(value='').sum(axis = 1).apply(lambda x : x.lower())
```

```
[23]: #get the final version of categorical, numerical, transactional and demographic
    ↪ cols

# update each col type with unwanted cols
cat_cols_final = cat_cols[:]
num_cols_final = num_cols[:]

t_cols_final = tagged_t_cols + t_cols[:-1]
d_cols_final = tagged_d_cols + d_cols[:-1]

for unwanted_col in unwanted_cols + form_change_cols + ['trans_datetime']:
    if unwanted_col in cat_cols:
        cat_cols_final.remove(unwanted_col)
    if unwanted_col in t_cols_final:
```

```

        t_cols_final.remove(unwanted_col)
    elif unwanted_col in d_cols_final:
        d_cols_final.remove(unwanted_col)
    else:
        pass
    elif unwanted_col in num_cols:
        num_cols_final.remove(unwanted_col)
    if unwanted_col in t_cols_final:
        t_cols_final.remove(unwanted_col)
    elif unwanted_col in d_cols_final:
        d_cols_final.remove(unwanted_col)
    else:
        pass
else:
    pass

# update each of them with new cols
t_cols_final.extend(['trans_year', 'trans_month'])
d_cols_final.extend(['age'])

cat_cols_final.extend(['trans_year', 'trans_month'])
num_cols_final.extend(['age'])

```

```

[32]: # inspecting the number of different types of columnsn
len(t_cols_final), len(d_cols_final), len(cat_cols_final), len(num_cols_final)

```

```

[32]: (17, 17, 27, 7)

```

```

[24]: # creating the preprocessed version of the dataset
eda_df = eda_df[d_cols_final + t_cols_final + ['is_fraud']]
print(f'Shape of the preprocessed dataset : {eda_df.shape}')

```

```

Shape of the preprocessed dataset : (3000, 35)

```

```

[25]: # TODO : remove this if required

eda_df.to_csv('/kaggle/working/eda_final_17_Apr_24.csv', index = False)

```

5 Data Visualization

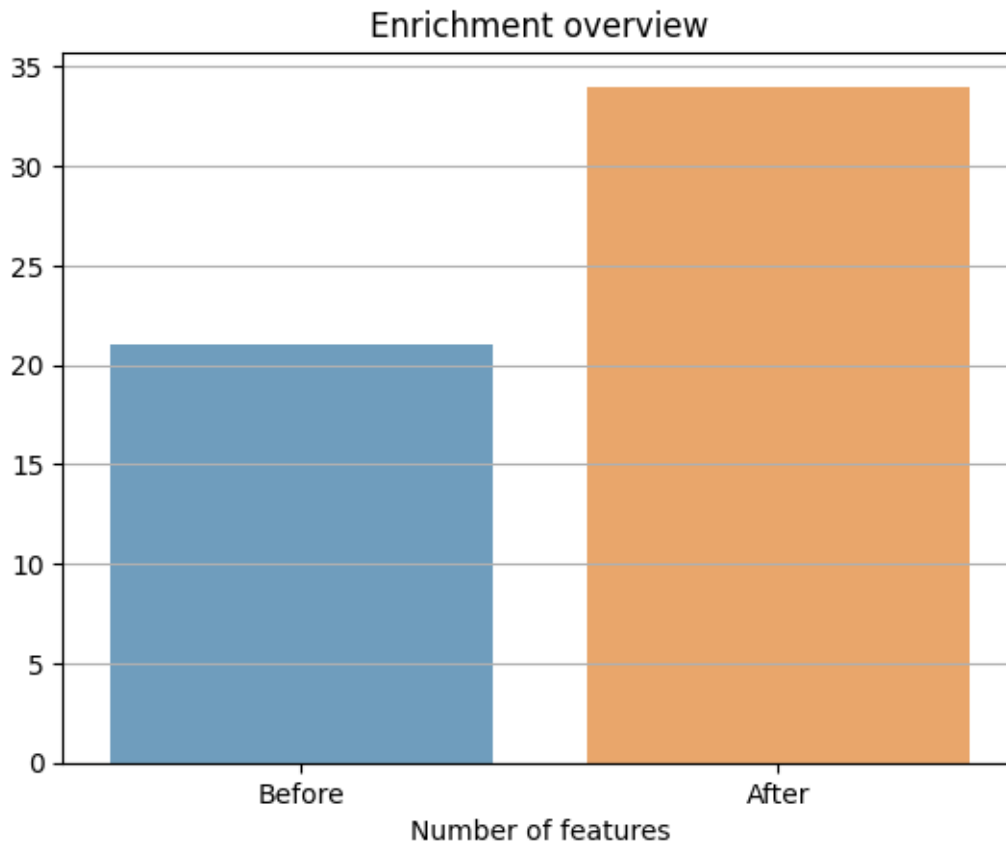
This is done for descriptive analytics as well as to find potential strategy for feature engineering and finally modelling the data

```

[29]: # Enrichment overview
n_features_0 = len(d_cols[:-1] + t_cols[:-1])
n_features_1 = len(d_cols_final + t_cols_final)

```

```
sns.barplot(x = ['Before', 'After'], y = [n_features_0, n_features_1], alpha = 0.7)
plt.grid(axis = 'y')
plt.xlabel('Number of features')
plt.title('Enrichment overview')
plt.show()
```



5.1 Demography based visualizations


5.1.1 Gender distribution for fraud

```
[30]: plt.figure(figsize = (6, 5))
ax = sns.countplot(data = eda_df, x = 'gender', hue = 'is_fraud',
                  alpha = 0.7,
                  # palette = 'dark',
                  )
male_total_count = eda_df[eda_df['gender'] == 'M']['gender'].count()
female_total_count = eda_df[eda_df['gender'] == 'F']['gender'].count()
counter = 0
```

```

for p in ax.patches:
    xcoord = p.get_x() + p.get_width()/2
    # if counter % 2 == 0:
    ycoord = p.get_height()
    barheight = int(100 * p.get_height()/male_total_count)
    counter += 1
    if counter % 2 == 0:
        barheight = int(100 * p.get_height()/female_total_count)
    ax.text(xcoord, ycoord, f'{barheight} %',
            fontsize=8, color='black', ha='center', va='bottom')
plt.title('Gender profile in fraudulent behaviour')
plt.legend(title = 'Fraudulent', labels = ['No', 'Yes'], loc = 'lower right')
plt.tight_layout()
plt.show()

```

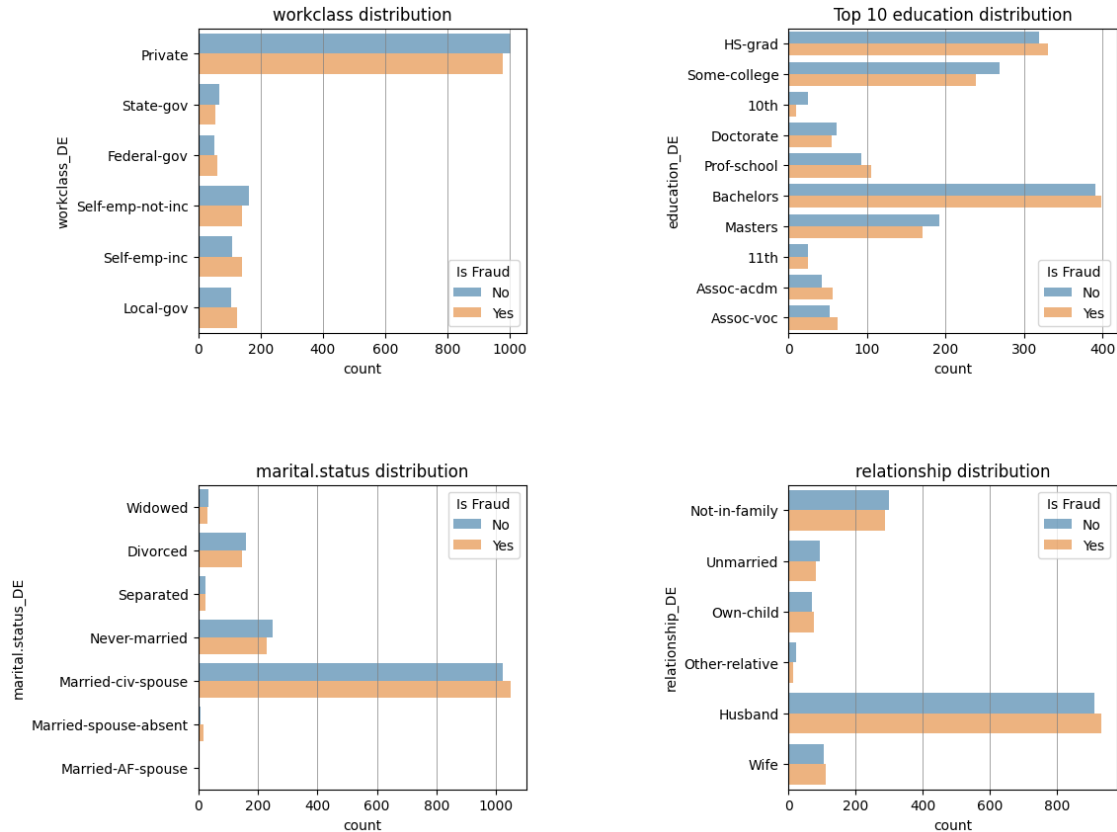


Observation : It can be seen that across genders, females are more related to Fraudulent transactions than males

```

[31]: fig, axes = plt.subplots(nrows = 2, ncols = 2, figsize = (12, 10))
plt.subplots_adjust(wspace = 0.8, hspace = 0.5)
for ax, xlabel in zip(axes.flatten(), d_cols_final[:4]):
    if eda_df[xlabel].unique().size <= 15 :
        ax.set_title(f'{xlabel.split("_")[0]} distribution')
        sns.countplot(eda_df, y = xlabel, hue = 'is_fraud', alpha = 0.6, ax = ax,
↪orient = 'h')
    elif (eda_df[xlabel].unique().size <= 100) and (type(eda_df[xlabel][0]) == int):
        ax.set_title(f'{xlabel.split("_")[0]} distribution')
        sns.histplot(eda_df,
                        y = xlabel,
                        hue = 'is_fraud',
                        multiple = 'dodge',
                        element = 'step',
                        alpha = 0.6,
                        ax = ax)
    else:
        top_k = 10
        ax.set_title(f'Top {top_k} {xlabel.split("_")[0]} distribution')
        freq_dist = eda_df[xlabel].value_counts(normalize = True) * 100
        top_entities = freq_dist[:top_k].index.tolist()
        sns.countplot(data = eda_df[eda_df[xlabel].isin(top_entities)],
                        y = xlabel, hue = 'is_fraud', ax = ax, orient = 'h', alpha = 0.6)
        ax.xaxis.grid(True, which = 'major', linestyle = '-', color = 'gray',
↪linewidth = 0.5)
        ax.legend(title = 'Is Fraud', labels = ['No', 'Yes'])
plt.show()

```

Observation (plots to be read in clockwise direction): 1. Among working class type, private job is heavily related to fraudulent transactions while state-government job is less related with it 2. Considering education background, high school graduates and those with bachelors degree are found to be more involved with fraud 3. Among the kind of relationship, the role “husband” is more associated with fraud 4. For the variable marital status, those married to civilian spouse are more associated with fraud than others

```
[32]: fig, axes = plt.subplots(nrows = 2, ncols = 2, figsize = (12, 10))
plt.subplots_adjust(wspace = 1.2, hspace = 0.5)
for ax, xlabel in zip(axes.flatten(), d_cols_final[4:8]):

    if eda_df[xlabel].unique().size <= 15 :
        ax.set_title(f'{xlabel.split("_")[0]} distribution')
        sns.countplot(eda_df, y = xlabel, hue = 'is_fraud', alpha = 0.6, ax = ax,
        ↪ax, orient = 'h')
    elif (eda_df[xlabel].unique().size <= 100) and (type(eda_df[xlabel][0]) ==_
    ↪int):
        ax.set_title(f'{xlabel.split("_")[0]} distribution')
        sns.histplot(eda_df,
                      y = xlabel,
                      hue = 'is_fraud',
```

```


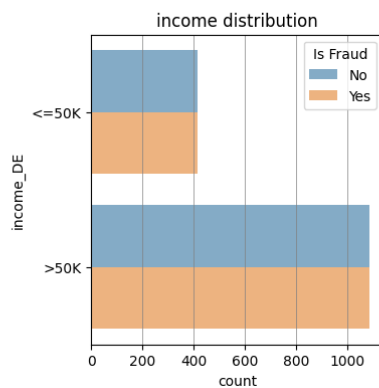
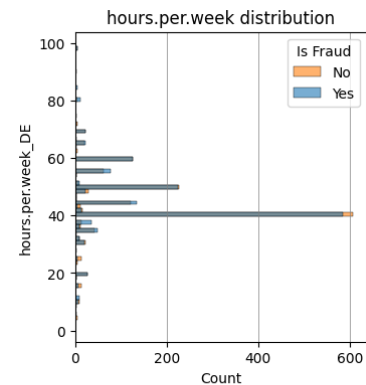

        multiple = 'layer',
        element = 'bars',
        alpha = 0.6,
        ax = ax)

else:
    top_k = 10
    ax.set_title(f'Top {top_k} {xlabel.split("_")[0]} distribution')
    freq_dist = eda_df[xlabel].value_counts(normalize = True) * 100
    top_entities = freq_dist[:top_k].index.tolist()
    sns.countplot(data = eda_df[eda_df[xlabel].isin(top_entities)],
                  y = xlabel,
                  hue = 'is_fraud',
                  ax = ax,
                  orient = 'h',
                  alpha = 0.6)

    ax.xaxis.grid(True, which = 'major', linestyle = '-', color = 'gray',
    ↪linewidth = 0.5)
    ax.legend(title = 'Is Fraud', labels = ['No', 'Yes'])

plt.show()

```



Observation (plots to be read in clockwise direction): 1. Race : Whites are more connected with fraud 2. Work hours : Those working about 40 hours a weeks appears to have more connection with fraud 3. Payment bank : Bank of America and Wells Fargo have highest fraud rates 4. Income distribution : Fraudulent activity is more prevalent people who earn more than \$50000 per annum

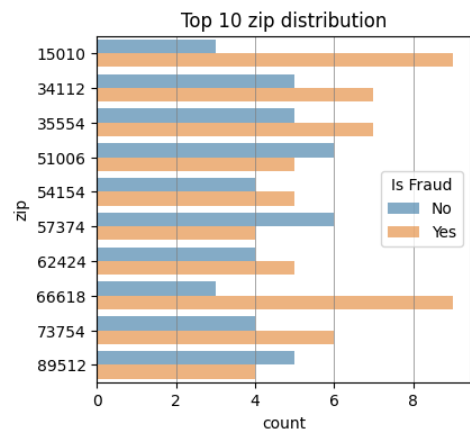
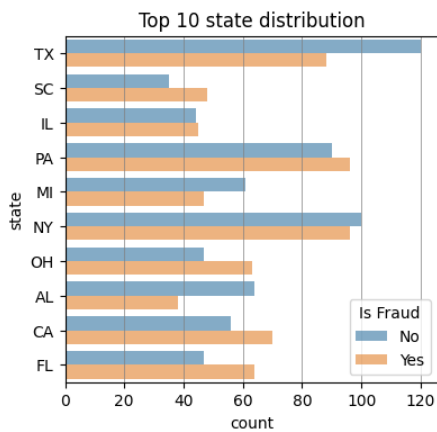
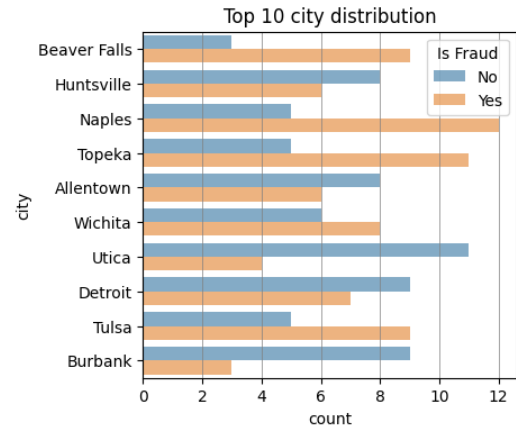
```
[33]: fig, axes = plt.subplots(nrows = 2, ncols = 2, figsize = (12, 10))
plt.subplots_adjust(wspace = 0.8, hspace = 0.5)
for ax, xlabel in zip(axes.flatten(), d_cols_final[8:12]):

    if eda_df[xlabel].unique().size <= 15 :
        ax.set_title(f'{xlabel.split("_")[0]} distribution')
        sns.countplot(eda_df, y = xlabel, hue = 'is_fraud', alpha = 0.6, ax = ax,
        ↪orient = 'h')
    elif (eda_df[xlabel].unique().size <= 100) and (type(eda_df[xlabel][0]) == int):
        ↪int):
        ax.set_title(f'{xlabel.split("_")[0]} distribution')
        sns.histplot(eda_df,
                        y = xlabel,
                        hue = 'is_fraud',
                        multiple = 'dodge',
                        element = 'step',
                        alpha = 0.6,
                        ax = ax)

    else:
        top_k = 10
        ax.set_title(f'Top {top_k} {xlabel.split("_")[0]} distribution')
        freq_dist = eda_df[xlabel].value_counts(normalize = True) * 100
        top_entities = freq_dist[:top_k].index.tolist()
        sns.countplot(data = eda_df[eda_df[xlabel].isin(top_entities)],
                        y = xlabel,
                        hue = 'is_fraud',
                        ax = ax,
                        orient = 'h',
                        alpha = 0.6)

        ax.xaxis.grid(True, which = 'major', linestyle = '-', color = 'gray',
        ↪linewidth = 0.5)
        ax.legend(title = 'Is Fraud', labels = ['No', 'Yes'])

plt.show()
```



Observation (plots to be read in clockwise direction): 1. Gender : Females are more connected with fraud 2. Customer city : Cities Naples and Topeka are have been 3. Customer zip code : This gives more granular view of the top 10 location of fraudulent activities 4. Customer state of residence : A high level view of fraudulent activity among top 10 states

```
[34]: fig, axes = plt.subplots(nrows = 3, ncols = 1, figsize = (8, 10))
plt.subplots_adjust(wspace = 0.8, hspace = 0.5)
for ax, xlabel in zip(axes.flatten(), d_cols_final[-3:]):

    if eda_df[xlabel].unique().size <= 100 :
        ax.set_title(f'{xlabel.split("_")[0]} distribution')
        sns.histplot(eda_df, y = xlabel, hue = 'is_fraud', alpha = 0.6, ax = ax)
    elif (eda_df[xlabel].unique().size <= 1000) and ((type(eda_df[xlabel][0]) ==
    int) or (type(eda_df[xlabel][0]) == np.int32) or (type(eda_df[xlabel][0]) ==
    np.int64)):
        unique_entries = eda_df[xlabel].unique().size
        binrange = None
        split_title = ' '.join(xlabel.split("_")[:2])
```


```

ax.set_title(f'{split_title} distribution')
sns.histplot(eda_df,
              y = xlabel,
              hue = 'is_fraud',
              multiple = 'layer',
              element = 'bars',
              alpha = 0.6,
              binrange = binrange,
              log_scale = True,
              ax = ax)

else:
    top_k = 10
    ax.set_title(f'Top {top_k} {xlabel} distribution')
    freq_dist = eda_df[xlabel].value_counts(normalize = True) * 100
    top_entities = freq_dist[:top_k].index.tolist()
    sns.countplot(data = eda_df[eda_df[xlabel].isin(top_entities)],
                  y = xlabel,
                  hue = 'is_fraud',
                  ax = ax,
                  orient = 'h',
                  alpha = 0.6)
    ax.xaxis.grid(True, which = 'major', linestyle = '-', color = 'gray',
                  linewidth = 0.5)
    ax.legend(title = 'Is Fraud', labels = ['No', 'Yes'])

plt.show()

```



Observation (plots to be read in top to bottom direction): 1. Customer city population : Fraudulent behaviour is more found in cities with population close to 1000 residents 2. Job designation of customer : Those performing as engineers and officer tend to be subjected to fraudulent activities more 3. Customer age : Customers in and around the age of 50 are subjected to fraudulent activities more

5.2 Transaction based visualizations

```
[36]: # 'type_TE' - removed as it does not corresponds to credit card transaction
      ↪ data
vis_t_cols = [
    'brand_TE',
    'Use Chip_TE',
    'MCC_TE',
    'shopping_mall_TE',
    'merchant',
    'category',
    'amt',
    'trans_year',
    'trans_month']
len(vis_t_cols)
```

[36]: 9

```
[39]: fig, axes = plt.subplots(nrows = 2, ncols = 2, figsize = (15,15))
      plt.subplots_adjust(wspace = 0.35, hspace = 0.4)
      for ax, pltlabel in zip(axes.flatten(), vis_t_cols[:4]):

          label_cardinality = eda_df[pltlabel].unique().size
          int_types = (int, np.int16, np.int32, np.int64)

          # print(f'{pltlabel} cardinality : ', label_cardinality)
          #if col is categorical <= 15 cardinality
          if (label_cardinality <= 15) and (not isinstance(eda_df[pltlabel][0],
          ↪int_types)):
              ax.set_title(f'{pltlabel.split("_")[0]} distribution')
              sns.countplot(eda_df,
                             y = pltlabel,
                             hue = 'is_fraud',
                             alpha = 0.6,
                             ax = ax,
                             orient = 'h')

          #if col is categorical with > 15 cardinality
          elif (label_cardinality > 15) and (not isinstance(eda_df[pltlabel][0],
          ↪int_types)):
              top_k = 10
              ax.set_title(f'Top {top_k} {pltlabel.split("_")[0]} distribution')
              freq_dist = eda_df[pltlabel].value_counts(normalize = True) * 100
              top_entities = freq_dist[:top_k].index.tolist()
              sns.countplot(data = eda_df[eda_df[pltlabel].isin(top_entities)],
                             y = pltlabel,
```

```

        hue = 'is_fraud',
        ax = ax,
        orient = 'h',
        alpha = 0.6)

#if col is numerical with cardinality <= 100 and element dtype is any int
    elif (label_cardinality <=100) and isinstance(eda_df[pltlabel][0],  

    ↪int_types):

        if 'MCC' in pltlabel:
            top_k = 10
            ax.set_title(f'Top {top_k} {pltlabel.split("_")[0]} distribution')
            freq_dist = eda_df[pltlabel].value_counts(normalize = True) * 100
            top_entities = freq_dist[:top_k].index.tolist()
            sns.countplot(data = eda_df[eda_df[pltlabel].isin(top_entities)],
                           y = pltlabel,
                           hue = 'is_fraud',
                           ax = ax,
                           orient = 'h',
                           alpha = 0.6)

        else:
            ax.set_title(f'{pltlabel.split("_")[0]} distribution')
            sns.histplot(eda_df,
                          y = pltlabel,
                          hue = 'is_fraud',
                          alpha = 0.6,
                          ax = ax)

    # if col is numerical with cardinativity <= 1000 and element dtype is any int
    elif (label_cardinality <= 3000) and isinstance(eda_df[pltlabel][0],  

    ↪int_types):
        ax.set_title(f'{pltlabel.split("_")[0]} distribution')
        sns.histplot(eda_df,
                      y = pltlabel,
                      hue = 'is_fraud',
                      alpha = 0.6,
                      log_scale = True,
                      ax = ax)

        ax.xaxis.grid(True, which = 'major', linestyle = '-', color = 'gray',  

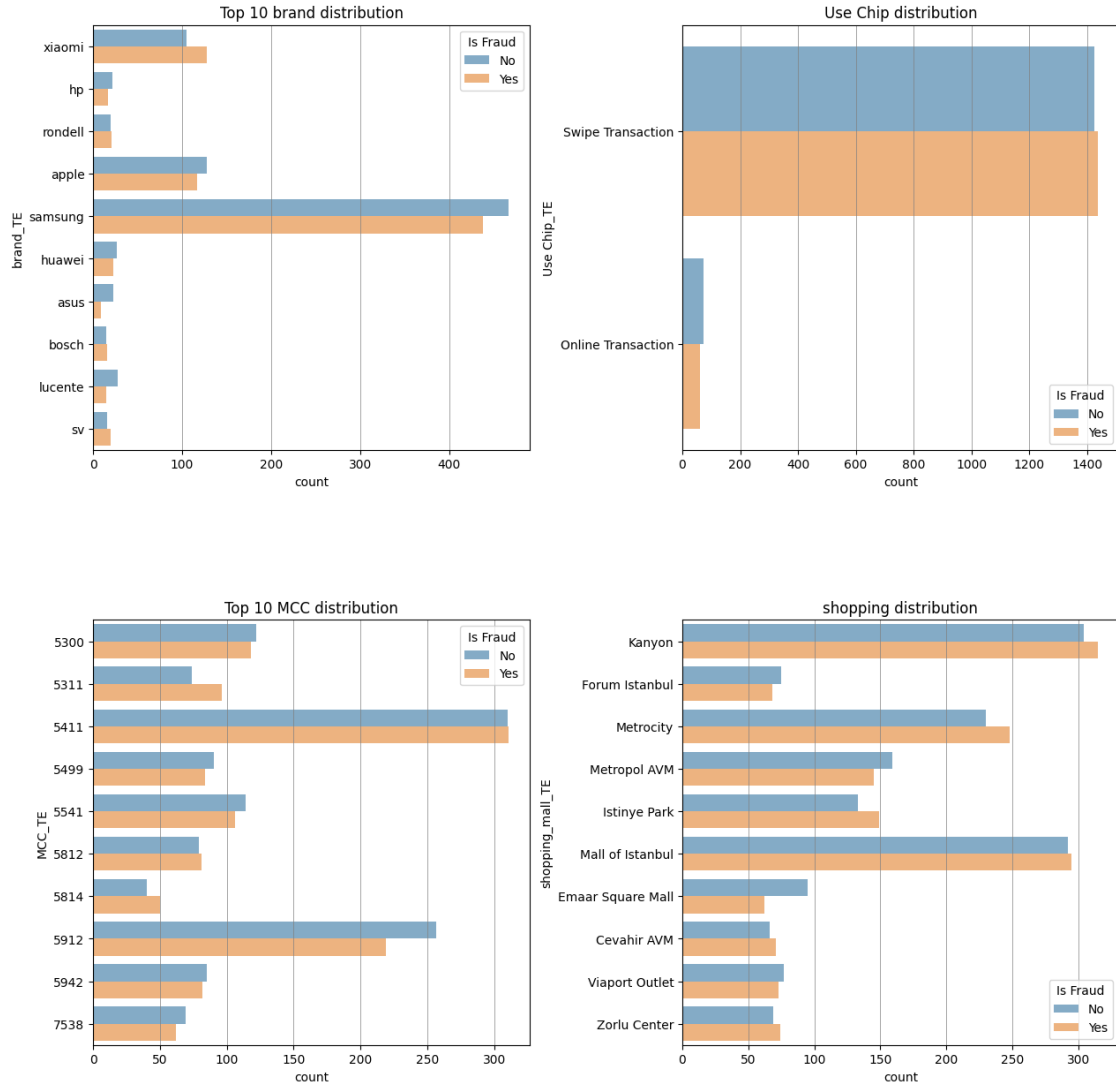
        ↪linewidth = 0.5)

        ax.legend(title = 'Is Fraud', labels = ['No', 'Yes'])#, loc = 'center',  

        ↪right', bbox_to_anchor = (-.3, 0.9)) #blue NO, Orange Yes

plt.show()

```

Observation (plots to be read in clockwise direction): 1. Product brand purchased : Fraudsters most commonly transact for samsung products 2. Transaction mode : Swipe mode of transaction has the most number of frauds reported 3. Purchase location : Shopping malls are the fraud hotspot with Mall of Istanbul at the top 4. Merchant Category Code (MCC) : 5411 (Grocery stores and supermarkets) is the dominant category subjected to fraudulent activities

```
[28]: fig, axes = plt.subplots(nrows = 2, ncols = 2, figsize = (15,15))
plt.subplots_adjust(wspace = 0.35, hspace = 0.4)
for ax, pltlabel in zip(axes.flatten(), vis_t_cols[4:8]):

    label_cardinality = eda_df[pltlabel].unique().size
    int_types = (int, np.int16, np.int32, np.int64)
    float_types = (float, np.float16, np.float32, np.float64)
```

```

#     print(f'{pltlabel} cardinality : ', label_cardinality)

# if col is float
if isinstance(eda_df[pltlabel][0], float_types):
    ax.set_title(f'{pltlabel.split("_")[0]} distribution')
    sns.histplot(eda_df,
                  y = pltlabel,
                  hue = 'is_fraud',
                  alpha = 0.5,
                  ax = ax)

# if col is categorical <= 15 cardinality
elif (label_cardinality <= 15) and (not isinstance(eda_df[pltlabel][0],
↳ int_types)):
    ax.set_title(f'{pltlabel.split("_")[0]} distribution')
    sns.countplot(eda_df,
                  y = pltlabel,
                  hue = 'is_fraud',
                  alpha = 0.6,
                  ax = ax,
                  orient = 'h')

# if col is categorical with > 15 cardinality
elif (label_cardinality > 15) and (not isinstance(eda_df[pltlabel][0],
↳ int_types)):
    top_k = 10
    ax.set_title(f'Top {top_k} {pltlabel.split("_")[0]} distribution')
    freq_dist = eda_df[pltlabel].value_counts(normalize = True) * 100
    top_entities = freq_dist[:top_k].index.tolist()
    sns.countplot(data = eda_df[eda_df[pltlabel].isin(top_entities)],
                  y = pltlabel,
                  hue = 'is_fraud',
                  ax = ax,
                  orient = 'h',
                  alpha = 0.6)

# if col is numerical with cardinality <= 100 and element dtype is any int
elif (label_cardinality <= 100) and isinstance(eda_df[pltlabel][0],
↳ int_types):

    if 'MCC' in pltlabel:
        top_k = 10
        ax.set_title(f'Top {top_k} {pltlabel.split("_")[0]} distribution')
        freq_dist = eda_df[pltlabel].value_counts(normalize = True) * 100

```

```

top_entities = freq_dist[:top_k].index.tolist()
sns.countplot(data = eda_df[eda_df[pltlabel].isin(top_entities)],
               y = pltlabel,
               hue = 'is_fraud',
               ax = ax,
               orient = 'h',
               alpha = 0.6)
elif label_cardinality < 10:
    split_title = ' '.join(pltlabel.split("_")[:2])
    ax.set_title(f'{split_title} distribution')
    sns.countplot(eda_df,
                  y = pltlabel,
                  hue = 'is_fraud',
                  alpha = 0.6,
                  ax = ax)

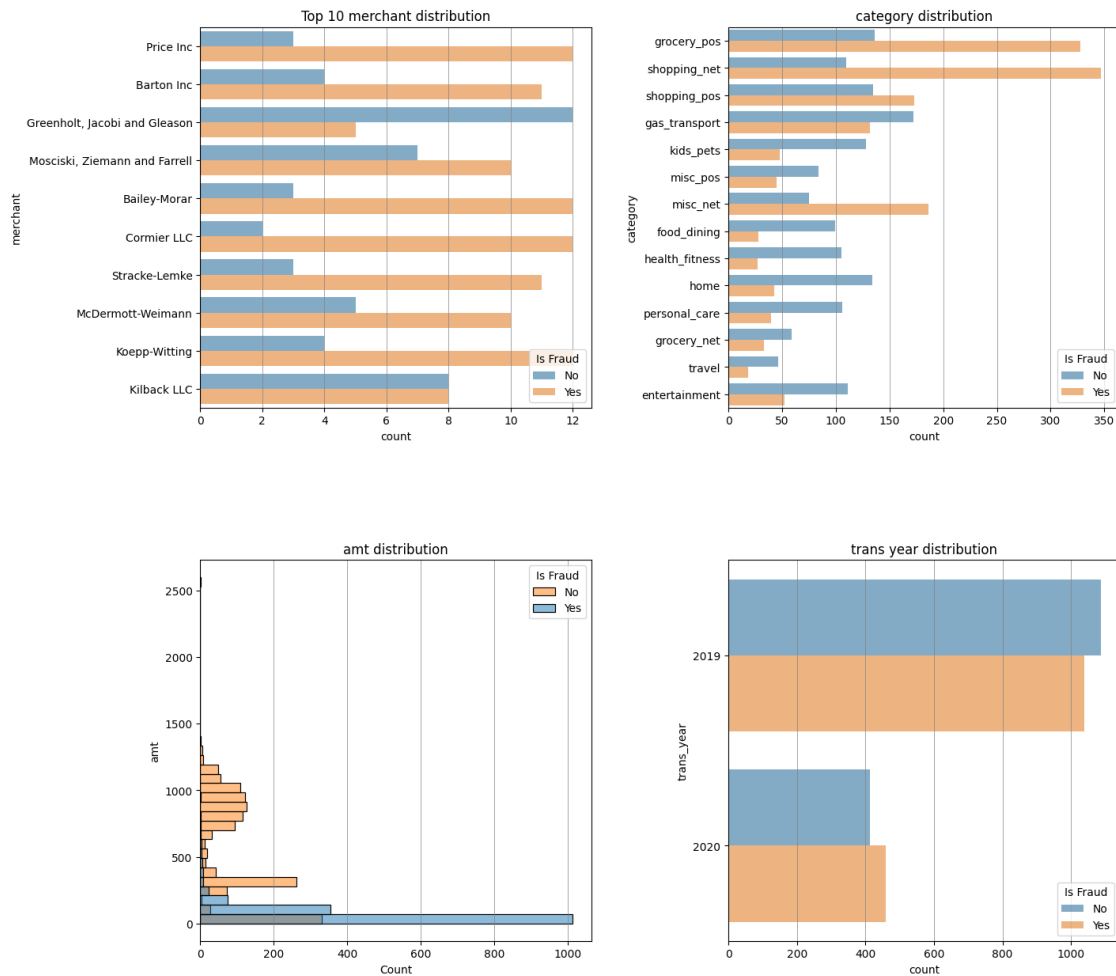
else:
    ax.set_title(f'{pltlabel.split("_")[0]} distribution')
    sns.histplot(eda_df,
                 y = pltlabel,
                 hue = 'is_fraud',
                 alpha = 0.6,
                 ax = ax)

# if col is numerical with cardinativity <= 1000 and element dtype is any int
elif (label_cardinality <= 3000) and isinstance(eda_df[pltlabel][0],
↪int_types):
    ax.set_title(f'{pltlabel.split("_")[0]} distribution')
    sns.histplot(eda_df,
                  y = pltlabel,
                  hue = 'is_fraud',
                  alpha = 0.6,
                  log_scale = True,
                  ax = ax)

    ax.xaxis.grid(True, which = 'major', linestyle = '-', color = 'gray',
↪linewidth = 0.5)
    ax.legend(title = 'Is Fraud', labels = ['No', 'Yes'])#, loc = 'center',
↪right', bbox_to_anchor = (-.3, 0.9)) #blue NO, Orange Yes

plt.show()

```



Observation (plots to be read in clockwise direction): 1. Merchant Name : Price Inc, Baily Morar, Cornier-LLC and Koepp-Witting are the top names associate with fraudulent activities 2. Shopping category : Online shopping and POS-sale of grocery at the top category of fraud 3. Transaction year : 2019 has more report of fraud than the year 2020 (pandemic time) 4. Transaction value : Fraudulent activity is focused more at large number of small value transactions typically averaging about \$250

```
[41]: pltlabel = vis_t_cols[-1]
fig = plt.figure(figsize = (6,6))
ax = plt.gca()

label_cardinality = eda_df[pltlabel].unique().size
int_types = (int, np.int16, np.int32, np.int64)
float_types = (float, np.float16, np.float32, np.float64)

# print(f'{pltlabel} cardinality : ', label_cardinality)
```

```

# if col is float
if isinstance(eda_df[pltlabel][0], float_types):
    ax.set_title(f'{pltlabel.split("_")[0]} distribution')
    sns.histplot(eda_df,
                  y = pltlabel,
                  hue = 'is_fraud',
                  alpha = 0.5,
                  ax = ax)

# if col is categorical <= 15 cardinality
elif (label_cardinality <= 15) and (not isinstance(eda_df[pltlabel][0], int_types)):
    split_title = ' '.join(pltlabel.split("_")[:2])
    ax.set_title(f'{split_title} distribution')
    sns.countplot(eda_df,
                  y = pltlabel,
                  hue = 'is_fraud',
                  alpha = 0.6,
                  ax = ax,
                  orient = 'h')

# if col is categorical with > 15 cardinality
elif (label_cardinality > 15) and (not isinstance(eda_df[pltlabel][0], int_types)):
    top_k = 10
    ax.set_title(f'Top {top_k} {pltlabel.split("_")[0]} distribution')
    freq_dist = eda_df[pltlabel].value_counts(normalize = True) * 100
    top_entities = freq_dist[:top_k].index.tolist()
    sns.countplot(data = eda_df[eda_df[pltlabel].isin(top_entities)],
                  y = pltlabel,
                  hue = 'is_fraud',
                  ax = ax,
                  orient = 'h',
                  alpha = 0.6)

# if col is numerical with cardinality <= 100 and element dtype is any int
elif (label_cardinality <= 100) and isinstance(eda_df[pltlabel][0], int_types):
    if 'MCC' in pltlabel:
        top_k = 10
        ax.set_title(f'Top {top_k} {pltlabel.split("_")[0]} distribution')
        freq_dist = eda_df[pltlabel].value_counts(normalize = True) * 100
        top_entities = freq_dist[:top_k].index.tolist()
        sns.countplot(data = eda_df[eda_df[pltlabel].isin(top_entities)],
                      y = pltlabel,
                      hue = 'is_fraud',

```

```


        ax = ax,
        orient = 'h',
        alpha = 0.6)
elif label_cardinality <= 15:
    split_title = ' '.join(pltlabel.split("_")[:2])
    ax.set_title(f'{split_title} distribution')
    sns.countplot(eda_df,
                  y = pltlabel,
                  hue = 'is_fraud',
                  alpha = 0.6,
                  ax = ax)

else:
    ax.set_title(f'{pltlabel.split("_")[0]} distribution')
    sns.histplot(eda_df,
                 y = pltlabel,
                 hue = 'is_fraud',
                 alpha = 0.6,
                 ax = ax)

# if col is numerical with cardinativity <= 1000 and element dtype is any int
elif (label_cardinality <= 3000) and isinstance(eda_df[pltlabel][0], int_types):
    ax.set_title(f'{pltlabel.split("_")[0]} distribution')
    sns.histplot(eda_df,
                 y = pltlabel,
                 hue = 'is_fraud',
                 alpha = 0.6,
                 log_scale = True,
                 ax = ax)
ax.xaxis.grid(True, which = 'major', linestyle = '-', color = 'gray', linewidth=
    ↳ 0.5)
ax.legend(title = 'Is Fraud', labels = ['No', 'Yes'])#, loc = 'center right',
    ↳ bbox_to_anchor = (-.3, 0.9)) #blue NO, Orange Yes

plt.show()

```



Observation (plots to be read from top to bottom direction): 1. Transaction month : March is the top month for fraud transactions

6 Preparing Master Dataset

6.1 Encode categorical variables

Converting the categorical values to integers to facilitate the ML modeling process

```
[29]: cat_col_encoding_transformer= ColumnTransformer([('encoder',
↳OrdinalEncoder(dtype = int), cat_cols_final)], remainder='passthrough')
cat_col_encoding_transformer
```

```
[29]: ColumnTransformer(remainder='passthrough',
transformers=[('encoder', OrdinalEncoder(dtype=<class 'int'>),
['gender', 'city', 'state', 'zip', 'job',
```

```
'cc_num', 'merchant', 'category', 'trans_num',
'workclass_DE', 'education_DE',
'marital.status_DE', 'relationship_DE',
'race_DE', 'hours.per.week_DE', 'income_DE',
'company_DE', 'product_id_TE', 'brand_TE',
'user_id_TE', 'user_session_TE', 'type_TE',
'Use Chip_TE', 'MCC_TE', 'shopping_mall_TE',
'trans_year', 'trans_month']]))
```

```
[30]: # final_master_dataset
encoded_master_cat = pd.DataFrame(cat_col_encoding_transformer.
    ↳fit_transform(eda_df[cat_cols_final]), columns = cat_cols_final)
final_master_dataset = encoded_master_cat.join(eda_df[num_cols_final +
    ↳['is_fraud']], how = 'inner')
final_master_dataset.head(2)
```

```
[30]:
```

	gender	city	state	zip	job	cc_num	merchant	category	trans_num	\
0	0	743	42	709	93	776	188	4	1341	
1	0	545	23	608	14	64	298	11	2650	

	workclass_DE	education_DE	marital.status_DE	relationship_DE	race_DE	\
0	2	11	6	1	4	
1	2	11	6	1	4	

	hours.per.week_DE	income_DE	company_DE	product_id_TE	brand_TE	\
0	32	0	341	22	357	
1	13	0	359	503	162	

	user_id_TE	user_session_TE	type_TE	Use Chip_TE	MCC_TE	\
0	195	229	3	1	28	
1	255	428	3	1	31	

	shopping_mall_TE	trans_year	trans_month	lat	long	city_pop	\
0	4	0	6	32.2768	-95.3031	144160	
1	2	0	10	38.9829	-93.9757	9512	

	amt	merch_lat	merch_long	age	is_fraud
0	347.88	32.063337	-94.562374	40	1
1	15.38	39.142095	-93.700393	52	0

7 Feature Engineering : Eliminating Less Informative features

Reducing the number of features to 60 % of the original one. During each step one most redundant feature gets eliminated until 60 % of them remains


```
[32]: feature_selector = RFE(DecisionTreeClassifier(), n_features_to_select= 0.6,
    ↪step = 1)
X, y = final_master_dataset[d_cols_final + t_cols_final],
    ↪final_master_dataset['is_fraud']
feature_selector.fit(X, y)
selected_feature_indices = np.where(feature_selector.ranking_ == 1)[0]
final_dataset_cols = final_master_dataset.columns[selected_feature_indices].
    ↪tolist()
print(final_dataset_cols)
```

```
['gender', 'job', 'workclass_DE', 'education_DE', 'marital.status_DE',
'race_DE', 'hours.per.week_DE', 'income_DE', 'company_DE', 'product_id_TE',
'brand_TE', 'user_id_TE', 'user_session_TE', 'shopping_mall_TE', 'trans_year',
'trans_month', 'lat', 'long', 'city_pop', 'amt']
```

```
[33]: column_filter_func = lambda x, y : [val for val in x if val in y]

d_cols_final = column_filter_func(final_dataset_cols, d_cols_final)
t_cols_final = column_filter_func(final_dataset_cols, t_cols_final)

cat_cols_final = column_filter_func(final_dataset_cols, cat_cols_final)
num_cols_final = column_filter_func(final_dataset_cols, num_cols_final)
```

```
[34]: # DATASET_DIR = 'datasets'
final_master_dataset = final_master_dataset[d_cols_final + t_cols_final +
    ↪['is_fraud']]
final_master_dataset.to_csv(f'{DATASET_DIR}/final_master_dataset.csv', index
    ↪=False)
```

```
[48]: final_master_dataset.head(2)
```

```
[48]:
```

	gender	state	workclass_DE	education_DE	marital.status_DE	race_DE	\
0	1	17	2	11	6	4	
1	0	42	2	11	6	4	

	hours.per.week_DE	income_DE	company_DE	lat	long	city_pop	age	\
0	32	0	341	30.0252	-90.2522	68211	90	
1	13	0	359	29.7972	-95.3288	2906700	41	

	product_id_TE	user_id_TE	user_session_TE	shopping_mall_TE	trans_year	\
0	22	195	229	4	1	
1	503	255	428	2	0	

	trans_month	amt	is_fraud
0	4	928.47	1
1	0	34.99	0

```
[49]: final_master_dataset.shape
```

```
[49]: (3000, 21)
```

8 Build HCCFD Model Variant 1 : Mixed Feature Model

This variant makes use of the dataset that contains both the transactional and demographic features of the master dataset

8.1 Partitioning dataset into train and test

Splitting ratio is train : test :: 80 : 20

```
[8]: #splitting dataset
predictors = d_cols_final + t_cols_final
target = ['is_fraud']
X, y = final_master_dataset[predictors], final_master_dataset[target]
test_split_ratio = 0.2
xtrain, xtest, ytrain, ytest = train_test_split(X, y, test_size = 0.2,
        ↪test_split_ratio, random_state = 10)
```

```
[10]: def get_baseline_accuracy(fitted_model, xtrain, xtest, ytrain, ytest):
        train_acc, test_acc = round(fitted_model.score(xtrain, ytrain) * 100, 2),
        ↪round(fitted_model.score(xtest, ytest) * 100, 2)
        return train_acc, test_acc
```

```
[11]: def show_accuracy(train_test_accuracy):
        train_acc, test_acc = train_test_accuracy
        print(f'Train Acc : {train_acc:3.2f} %', end = ' || ')
        print(f'Test Acc : {test_acc:3.2f} %')
```

##

Decision Tree

8.2 Baseline Model

```
[50]: # baseline model
# dt
dt = DecisionTreeClassifier(random_state = 10)
dt.fit(xtrain, ytrain)

train_preds = dt.predict(xtrain)
test_preds = dt.predict(xtest)

dt_baseline_accuracy = get_baseline_accuracy(dt, xtrain, xtest, ytrain, ytest)
show_accuracy(dt_baseline_accuracy)
```

Train Acc : 100.00 % || Test Acc : 84.50 %

8.3 Probabilistic Analysis

```
[12]: TARGET_LABELS = {0 : 'Not Fraud', 1 : 'Fraud'}
def show_classification_report(ytrue, ypred):
    report= classification_report(y_true= ytrue, y_pred= ypred, output_dict = True)
    final_report = {}
    for report_key, report_value in report.items():
        if report_key.isnumeric():
            if int(report_key) in list(TARGET_LABELS.keys()):
                final_report[TARGET_LABELS[int(report_key)]] = report_value
    return pd.DataFrame(final_report)
```

```
[13]: def show_roc_auc_score(ytrain, trainpreds, ytest, testpreds):
    print(f'Train ROC-AUC score : {roc_auc_score(ytrain, trainpreds):.3f} ', end = ' || ')
    print(f'Test ROC-AUC score : {roc_auc_score(ytest, testpreds) :.3f}')
```

```
[14]: def show_cv_score(fitted_model, xtrain, ytrain):
    cv_score = cross_val_score(fitted_model, xtrain, ytrain).mean() * 100
    print(f'Train Cross validation Acc : {cv_score :.3f} %')
```

```
[51]: dt_train_report = show_classification_report(ytrain, train_preds)
dt_train_report
```

```
[51]:
```

	Not Fraud	Fraud
precision	1.0	1.0
recall	1.0	1.0
f1-score	1.0	1.0
support	1215.0	1185.0

```
[52]: dt_test_report = show_classification_report(ytest, test_preds)
dt_test_report
```

```
[52]:
```

	Not Fraud	Fraud
precision	0.817881	0.872483
recall	0.866667	0.825397
f1-score	0.841567	0.848287
support	285.000000	315.000000

```
[53]: show_roc_auc_score(ytrain, train_preds, ytest, test_preds)
```

Train ROC-AUC score : 1.000 || Test ROC-AUC score : 0.846

```
[54]: show_cv_score(dt, xtrain, ytrain)
```

Train Cross validation Acc : 83.000 %

```
[55]: # save trained model
joblib.dump(dt, os.path.join(MODEL_DIR, 'baseline_dt.joblib'))
```

```
[55]: ['./models/baseline_dt.joblib']
```

8.4 HyperParameter Tuning

```
[15]: def test_model_with_best_params(model_obj, best_params):
    model = model_obj(**best_params)
    model.fit(xtrain, ytrain)

    train_acc_score, test_acc_score = round(model.score(xtrain, ytrain) * 100, 2), round(model.score(xtest, ytest) * 100, 2)
    return train_acc_score, test_acc_score, model
```

```
[103]: def dt_objective_fn(params_to_tune):
    decision_tree = DecisionTreeClassifier(**params_to_tune)
    accuracy_score = cross_val_score(decision_tree, xtrain, ytrain, cv = 5).mean()
    return {'loss': -accuracy_score, 'status': STATUS_OK}

# hyperparameter search space as python dict
tune_params = ['criterion', 'min_samples_split']
search_space = {
    'criterion': hp.choice('criterion', ["gini", "entropy", "log_loss"]),
    'min_samples_split': hp.uniform('min_samples_split', 0.01, 1.0)
}

trials = Trials()
best_params = fmin(fn = dt_objective_fn,
                  space = search_space,
                  algo = tpe.suggest,
                  max_evals = 25,
                  trials = trials,
                  rstate = np.random.seed(10))

best_params
```

```
100%|          | 25/25 [00:02<00:00, 11.39trial/s, best loss:
-0.8775000000000001]
```

```
[103]: {'criterion': 0, 'min_samples_split': 0.2702447741874797}
```

```
[57]: best_params['criterion'] = ["gini", "entropy",
    ↪ "log_loss"][best_params['criterion']]
```

```
best_params
```

```
[57]: {'criterion': 'gini', 'min_samples_split': 0.2702447741874797}
```

```
[58]: dt_tuned_artifacts = test_model_with_best_params(DecisionTreeClassifier,  
↳ best_params)  
dt_tuned_accuracy, dt_tuned = dt_tuned_artifacts[:2], dt_tuned_artifacts[-1]  
show_accuracy(dt_tuned_accuracy)
```

Train Acc : 87.62 % || Test Acc : 88.17 %

```
[59]: train_preds = dt_tuned.predict(xtrain)  
test_preds = dt_tuned.predict(xtest)
```

```
[60]: dt_tuned_train_report = show_classification_report(ytrain, train_preds)  
dt_tuned_train_report
```

```
[60]:
```

	Not Fraud	Fraud
precision	0.948242	0.822674
recall	0.799177	0.955274
f1-score	0.867351	0.884030
support	1215.000000	1185.000000

```
[61]: dt_tuned_test_report = show_classification_report(ytest, test_preds)  
dt_tuned_test_report
```

```
[61]:
```

	Not Fraud	Fraud
precision	0.914729	0.856725
recall	0.828070	0.930159
f1-score	0.869245	0.891933
support	285.000000	315.000000

```
[62]: show_roc_auc_score(ytrain, train_preds, ytest, test_preds)
```

Train ROC-AUC score : 0.877 || Test ROC-AUC score : 0.879

```
[63]: show_cv_score(dt_tuned, xtrain, ytrain)
```

Train Cross validation Acc : 87.167 %

```
[111]: # save trained model  
joblib.dump(dt_tuned, os.path.join(MODEL_DIR, 'tuned_dt.joblib'))
```

```
[111]: ['/kaggle/working/models/tuned_dt.joblib']
```

```
[64]: def save_model(model_name, fitted_model):  
joblib.dump(fitted_model, os.path.join(MODEL_DIR, f'{model_name}.joblib'))
```

```
print(f'Model saved to {os.path.join(MODEL_DIR, f"{model_name}.joblib")}'))
```

```
[65]: def save_tuned_params(params:dict , model_name):  
      filepath = os.path.join(MODEL_DIR, f'tuned_{model_name}_params.json')  
      with open(filepath, 'w') as file:  
          json.dump(params, file)  
      print(f'File written to {filepath}')
```

```
[66]: save_tuned_params(best_params, 'dt')
```

File written to ./models/tuned_dt_params.json

##

Naive Bayes

8.5 Baseline Model

```
[67]: # nb  
naivebayes = GaussianNB()  
naivebayes.fit(xtrain, ytrain)  
  
train_preds = naivebayes.predict(xtrain)  
test_preds = naivebayes.predict(xtest)  
  
nb_baseline_accuracy = get_baseline_accuracy(naivebayes, xtrain, xtest, ytrain,   
↪ ytest)  
show_accuracy(nb_baseline_accuracy)
```

Train Acc : 85.25 % || Test Acc : 84.17 %

8.6 Probabilitic Analysis

```
[68]: nb_train_report = show_classification_report(ytrain, train_preds)  
nb_train_report
```

```
[68]:
```

	Not Fraud	Fraud
precision	0.788346	0.958104
recall	0.968724	0.733333
f1-score	0.869276	0.830784
support	1215.000000	1185.000000

```
[69]: nb_test_report = show_classification_report(ytest, test_preds)  
nb_test_report
```

```
[69]:
```

	Not Fraud	Fraud
precision	0.758152	0.974138
recall	0.978947	0.717460

```
f1-score      0.854518    0.826325
support      285.000000   315.000000
```

```
[70]: show_roc_auc_score(ytrain, train_preds, ytest, test_preds)
```

```
Train ROC-AUC score : 0.851 || Test ROC-AUC score : 0.848
```

```
[71]: show_cv_score(naivebayes, xtrain, ytrain)
```

```
Train Cross validation Acc : 85.083 %
```

```
[72]: #save model
      save_model('baseline_nb', naivebayes)
```

```
Model saved to ./models/baseline_nb.joblib
```

8.7 Naive Bayes : HyperParameter Tuning

```
[121]: def nb_objective_fn(params_to_tune):
        model = GaussianNB(**params_to_tune)
        accuracy_score = cross_val_score(model, xtrain, ytrain, cv = 5).mean()
        return {'loss': -accuracy_score , 'status' : STATUS_OK}

        # hyperparameter search space as python dict
        # tune_params = ['var_smoothing']
        search_space = {
            'var_smoothing' : hp.uniform('var_smoothing', 1e-20, 1e-6)
        }

        trials = Trials()
        best_params = fmin(fn = nb_objective_fn,
                           space = search_space,
                           algo = tpe.suggest,
                           max_evals = 25,
                           trials = trials,
                           rstate = np.random.seed(10))

        best_params
```

```
100%|      | 25/25 [00:01<00:00, 17.03trial/s, best loss:
-0.8591666666666666]
```

```
[121]: {'var_smoothing': 4.409409693293421e-09}
```

```
[74]: nb_tuned_artifacts = test_model_with_best_params(GaussianNB, best_params)
      nb_tuned_accuracy, nb_tuned = nb_tuned_artifacts[:2], nb_tuned_artifacts[-1]
      show_accuracy(nb_tuned_accuracy)
```

```
Train Acc : 85.25 % || Test Acc : 84.33 %
```

```
[75]: train_preds = nb_tuned.predict(xtrain)
      test_preds = nb_tuned.predict(xtest)
```

```
[76]: nb_tuned_train_report = show_classification_report(ytrain, train_preds)
      nb_tuned_train_report
```

```
[76]:
```

	Not Fraud	Fraud
precision	0.788346	0.958104
recall	0.968724	0.733333
f1-score	0.869276	0.830784
support	1215.000000	1185.000000

```
[77]: nb_tuned_test_report = show_classification_report(ytest, test_preds)
      nb_tuned_test_report
```

```
[77]:
```

	Not Fraud	Fraud
precision	0.760218	0.974249
recall	0.978947	0.720635
f1-score	0.855828	0.828467
support	285.000000	315.000000

```
[78]: show_roc_auc_score(ytrain, train_preds, ytest, test_preds)
```

Train ROC-AUC score : 0.851 || Test ROC-AUC score : 0.850

```
[79]: show_cv_score(nb_tuned, xtrain, ytrain)
```

Train Cross validation Acc : 85.167 %

```
[80]: #save model
      save_model('tuned_nb', nb_tuned)
```

Model saved to ./models/tuned_nb.joblib

```
[81]: #save params
      save_tuned_params(best_params, 'nb')
```

File written to ./models/tuned_nb_params.json

###

SVM

8.8 Baseline Model

```
[82]: # svm
      svm_cls = SVC(random_state=10)
      svm_cls.fit(xtrain, ytrain)
```



```

train_preds = svm_cls.predict(xtrain)
test_preds = svm_cls.predict(xtest)

svm_baseline_accuracy = get_baseline_accuracy(svm_cls, xtrain, xtest, ytrain,
↪ytest)
show_accuracy(svm_baseline_accuracy)

```

Train Acc : 52.12 % || Test Acc : 50.67 %

```

[83]: svm_train_report = show_classification_report(ytrain, train_preds)
      svm_train_report

```

```

[83]:
           Not Fraud      Fraud
precision    0.515820    0.557325
recall       0.885597    0.147679
f1-score     0.651924    0.233489
support     1215.000000  1185.000000

```

```

[84]: svm_test_report = show_classification_report(ytest, test_preds)
      svm_test_report

```

```

[84]:
           Not Fraud      Fraud
precision    0.489564    0.630137
recall       0.905263    0.146032
f1-score     0.635468    0.237113
support     285.000000  315.000000

```

```

[85]: show_roc_auc_score(ytrain, train_preds, ytest, test_preds)

```

Train ROC-AUC score : 0.517 || Test ROC-AUC score : 0.526

```

[86]: show_cv_score(svm_cls, xtrain, ytrain)

```

Train Cross validation Acc : 51.458 %

```

[87]: #save model
      save_model('baseline_svm', svm_cls)

```

Model saved to ./models/baseline_svm.joblib

8.9 HyperParameter Tuning

```

[81]: def svm_objective_fn(params_to_tune):
      model = SVC(**params_to_tune)
      accuracy_score = cross_val_score(model, xtrain, ytrain, cv = 5).mean()
      return {'loss': -accuracy_score , 'status' : STATUS_OK}

      # hyperparameter search space as python dict

```

```

search_space = {
    'C' : hp.lognormal('C', 0, 10),
    # 'kernel' : hp.choice('kernel_choices', ['linear', 'sigmoid'])
}

trials = Trials()
best_params = fmin(fn = svm_objective_fn,
                   space = search_space,
                   algo = tpe.suggest,
                   max_evals = 10,
                   trials = trials,
                   rstate = np.random.seed(10))

best_params

```

```

100%|          | 10/10 [03:48<00:00, 22.84s/trial, best loss:
-0.8429166666666668]

```

```
[81]: {'C': 2732960.906264294}
```

```
[88]: svm_tuned_artifacts = test_model_with_best_params(SVC, best_params)
      svm_tuned_accuracy, svm_tuned = svm_tuned_artifacts[:2], svm_tuned_artifacts[-1]
      show_accuracy(svm_tuned_accuracy)

```

```
Train Acc : 86.21 % || Test Acc : 85.50 %
```

```
[89]: train_preds = svm_tuned.predict(xtrain)
      test_preds = svm_tuned.predict(xtest)

```

```
[90]: svm_tuned_train_report = show_classification_report(ytrain, train_preds)
      svm_tuned_train_report

```

```
[90]:
```

	Not Fraud	Fraud
precision	0.808229	0.942029
recall	0.953909	0.767932
f1-score	0.875047	0.846118
support	1215.000000	1185.000000

```
[91]: svm_tuned_test_report = show_classification_report(ytest, test_preds)
      svm_tuned_test_report

```

```
[91]:
```

	Not Fraud	Fraud
precision	0.782857	0.956000
recall	0.961404	0.758730
f1-score	0.862992	0.846018
support	285.000000	315.000000

```
[92]: show_roc_auc_score(ytrain, train_preds, ytest, test_preds)
```

Train ROC-AUC score : 0.861 || Test ROC-AUC score : 0.860

```
[93]: show_cv_score(svm_tuned, xtrain, ytrain)
```

Train Cross validation Acc : 84.875 %

```
[94]: #save model
save_model('tuned_svm', svm_tuned)
```

Model saved to ./models/tuned_svm.joblib

```
[66]: #save model
save_tuned_params(best_params, 'svm')
```

File written to models/tuned_svm_params.json

###

Logistic Regression

8.10 Baseline Model

```
[95]: # baseline model
# dt
lr = LogisticRegression(random_state = 10)
lr.fit(xtrain, ytrain)

train_preds = lr.predict(xtrain)
test_preds = lr.predict(xtest)

lr_baseline_accuracy = get_baseline_accuracy(lr, xtrain, xtest, ytrain, ytest)
show_accuracy(lr_baseline_accuracy)
```

Train Acc : 84.83 % || Test Acc : 84.50 %

```
[96]: # baseline model
# dt
lr = LogisticRegression(random_state = 10)
lr.fit(xtrain, ytrain)

train_preds = lr.predict(xtrain)
test_preds = lr.predict(xtest)

lr_baseline_accuracy = get_baseline_accuracy(lr, xtrain, xtest, ytrain, ytest)
show_accuracy(lr_baseline_accuracy)
```

Train Acc : 84.83 % || Test Acc : 84.50 %

8.11 Probabilistic Analysis

```
[97]: lr_train_report = show_classification_report(ytrain, train_preds)
lr_train_report
```

```
[97]:
```

	Not Fraud	Fraud
precision	0.798178	0.921891
recall	0.937449	0.756962
f1-score	0.862226	0.831325
support	1215.000000	1185.000000

```
[98]: lr_test_report = show_classification_report(ytest, test_preds)
lr_test_report
```

```
[98]:
```

	Not Fraud	Fraud
precision	0.774286	0.944000
recall	0.950877	0.749206
f1-score	0.853543	0.835398
support	285.000000	315.000000

```
[99]: show_roc_auc_score(ytrain, train_preds, ytest, test_preds)
```

Train ROC-AUC score : 0.847 || Test ROC-AUC score : 0.850

```
[100]: show_cv_score(lr, xtrain, ytrain)
```

Train Cross validation Acc : 84.958 %

```
[102]: # save trained model
save_model('baseline_lr', lr)
```

Model saved to ./models/baseline_lr.joblib

8.12 HyperParameter Tuning

```
[25]: def lr_objective_fn(params_to_tune):
    logistic_reg = LogisticRegression(**params_to_tune)
    accuracy_score = cross_val_score(logistic_reg, xtrain, ytrain, cv = 5).
    ↪mean()
    return {'loss': -accuracy_score , 'status' : STATUS_OK}

# hyperparameter search space as python dict
search_space = {
    'solver' : hp.choice('solver', ['lbfgs', 'liblinear', 'newton-cg',
    ↪'newton-cholesky', 'sag', 'saga']),
    'C' : hp.lognormal('C', 0, 1)
}
```

```
# trials = Trials()
best_params = fmin(fn = lr_objective_fn,
                   space = search_space,
                   algo = tpe.suggest,
                   max_evals = 25,
                   rstate = np.random.seed(10)
                   )
best_params
```

100%| | 25/25 [00:09<00:00, 2.65trial/s, best loss: -0.85875]

[25]: {'C': 0.3575637936038211, 'solver': 2}

```
[26]: best_params['solver'] = ['lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky',
    ↪ 'sag', 'saga'][best_params['solver']]
best_params
```

[26]: {'C': 0.3575637936038211, 'solver': 'newton-cg'}

```
[104]: lr_tuned_artifacts = test_model_with_best_params(LogisticRegression,
    ↪ best_params)
lr_tuned_accuracy, lr_tuned = lr_tuned_artifacts[:2], lr_tuned_artifacts[-1]
show_accuracy(lr_tuned_accuracy)
```

Train Acc : 85.62 % || Test Acc : 85.50 %

```
[105]: train_preds = lr_tuned.predict(xtrain)
test_preds = lr_tuned.predict(xtest)
```

```
[106]: lr_tuned_train_report = show_classification_report(ytrain, train_preds)
lr_tuned_train_report
```

```
[106]:
```

	Not Fraud	Fraud
precision	0.802083	0.937500
recall	0.950617	0.759494
f1-score	0.870056	0.839161
support	1215.000000	1185.000000

```
[107]: lr_tuned_test_report = show_classification_report(ytest, test_preds)
lr_tuned_test_report
```

```
[107]:
```

	Not Fraud	Fraud
precision	0.782857	0.956000
recall	0.961404	0.758730
f1-score	0.862992	0.846018
support	285.000000	315.000000

```
[108]: show_roc_auc_score(ytrain, train_preds, ytest, test_preds)
```

Train ROC-AUC score : 0.855 || Test ROC-AUC score : 0.860

```
[109]: show_cv_score(lr_tuned, xtrain, ytrain)
```

Train Cross validation Acc : 85.875 %

```
[110]: # save trained model
save_model('tuned_lr', lr_tuned)
```

Model saved to ./models/tuned_lr.joblib

```
[111]: save_tuned_params(best_params, 'lr')
```

File written to ./models/tuned_lr_params.json

###

Model Performance Summary

```
[112]: def compile_variant_results(baseline_tuned_metric_pairs: Tuple[Tuple[Tuple,
↳ Tuple], ...],
                                model_names: list,
                                eval_columns: list = ['baseline_train_acc',
↳ 'baseline_test_acc', 'tuned_train_acc', 'tuned_test_acc'],
                                only_baseline: bool = False):
    if only_baseline:
        baseline_tuned_pairs = baseline_tuned_metric_pairs
    else:
        baseline_tuned_pairs = []
        for bline, ftune in baseline_tuned_metric_pairs:
            baseline_tuned_pairs.append(np.concatenate((bline, ftune), axis =
↳ 0))
        baseline_tuned_pairs = tuple(baseline_tuned_pairs)

    variant_results = np.vstack(baseline_tuned_pairs)
    if only_baseline :
        eval_columns = eval_columns[:2]
    return pd.DataFrame(variant_results, columns= eval_columns, index =
↳ model_names)
```

```
[46]: def get_model_acc(modelname):
    #load baselines
    baseline = joblib.load(f'./models/baseline_{modelname}_v2.joblib')
    #get baseline acc
    baseline_acc = get_baseline_accuracy(baseline, xtrain, xtest, ytrain, ytest)
```

```

#load tuned
with open(f'./models/tuned_{modelname}_params_v2.json') as paramfile:
    bestparams = json.load(paramfile)

if modelname == 'dt':
    model = DecisionTreeClassifier
elif modelname == 'nb':
    model = GaussianNB
elif modelname == 'svm':
    model = SVC
elif modelname == 'lr':
    model = LogisticRegression

#get tuned acc
tuned_acc = test_model_with_best_params(model, bestparams)[:2]

return baseline_acc, tuned_acc

def get_all_model_acc(modelnames):
    accs = []
    for modelname in modelnames:
        accs.append(get_model_acc(modelname))
    return tuple(accs)

```

```

[113]: # results overview for Model Variant 1
metric_tuples = (
    (dt_baseline_accuracy, dt_tuned_accuracy),
    (nb_baseline_accuracy, nb_tuned_accuracy),
    (svm_baseline_accuracy, svm_tuned_accuracy),
    (lr_baseline_accuracy, lr_tuned_accuracy)
)
variant_1_results = compile_variant_results(metric_tuples, ['dt', 'nb', 'svm', 'lr'])
variant_1_results

```

```

[113]:

```

	baseline_train_acc	baseline_test_acc	tuned_train_acc	tuned_test_acc
dt	100.00	84.50	87.62	88.17
nb	85.25	84.17	85.25	84.33
svm	52.12	50.67	86.21	85.50
lr	84.83	84.50	85.62	85.50

```

[115]: # saving variant 1 result
variant_1_results.to_csv(f'{REPORTS_DIR}/
    variant_1_baseline_vs_tuned_performance.csv', index = False)

```

9 Recommendation

The best model out of the 8 kinds built is selected based on the tuned accuracy. The best model thus selected is Decision Tree with $\sim 88\%$ test accuracy upon tuning.

10 Build Model Variant 2 - Pipelined model With Concesus-driven Inference

This architecture uses 2 models, 1 trained on master transactional dataset and the other on master demographic dataset. Both models will produce prediction for each datapoint and the transaction is flagged as fraud or not only when both of them agrees on it and if there is any discrepancy it is counted as “not fraud”.

10.1 Pipelined Model Definition

A python class to construct the variant 2 model

```
[116]: class ClassifierTransformer(BaseEstimator, TransformerMixin):
    def __init__(self, classifier):
        self.classifier = classifier

    def fit(self, X, y=None):
        self.classifier.fit(X, y)
        return self

    def transform(self, X, y=None):
        return self.classifier.predict(X).reshape(-1, 1)

class PipelinedModel(BaseEstimator):

    def __init__(self, model, model_feature_names : List[List, ], model_names:
↳List):
        cls_transformers : List[Tuple] = []
        models = [model] * len(model_names)
        self.model = model
        self.model_names = model_names
        self.model_feature_names = model_feature_names
        for model, model_feature_set, model_name in zip(models,
↳model_feature_names, model_names):
            cls_transformers.append((model_name, ClassifierTransformer(model),
↳model_feature_set))
        self.fitted_model = ColumnTransformer(cls_transformers, remainder =
↳'passthrough')

    def fit(self, X, y):
        self.fitted_model.fit(X, y)
```



```

def predict(self, X):
    self.predictions = self.fitted_model.transform(X)
    hsplit_preds = np.hsplit(self.predictions, self.predictions.shape[1])
    final_preds = np.logical_and(*hsplit_preds).astype(int).reshape(-1,1)
    return final_preds

def get_baseline_accuracy(self, ytrue, ypred):
    acc = round(accuracy_score(y_true=ytrue, y_pred = ypred) * 100, 2)
    return acc

def score(self, X, y):
    predictions = self.predict(X)
    return accuracy_score(y_true = y, y_pred = predictions)

```

11 Pipelined Decision Tree Model with Consensus

```

[117]: base_model = DecisionTreeClassifier(random_state=10)
p_feature_names = [d_cols_final, t_cols_final]
p_model_names = ['d_predictor', 't_predictor']
p_dt_model = PipelinedModel(base_model, p_feature_names, p_model_names)

p_dt_model.fit(xtrain, ytrain)#.fit(xtrain, ytrain)

trainpreds = p_dt_model.predict(xtrain)
testpreds = p_dt_model.predict(xtest)

baseline_acc = namedtuple('Baseline_Accuracy', ['training', 'test'])

p_dt_baseline_acc = baseline_acc(training = p_dt_model.
    ↪get_baseline_accuracy(ytrain, trainpreds),
                                test = p_dt_model.get_baseline_accuracy(ytest,
    ↪testpreds))
show_accuracy(p_dt_baseline_acc)

```

Train Acc : 100.00 % || Test Acc : 67.17 %

11.1 Probabilistic Analysis

```

[118]: p_dt_train_report = show_classification_report(ytrain, trainpreds)
p_dt_train_report

```

```

[118]:

```

	Not Fraud	Fraud
precision	1.0	1.0
recall	1.0	1.0
f1-score	1.0	1.0
support	1215.0	1185.0

```
[119]: p_dt_test_report = show_classification_report(ytest, testpreds)
p_dt_test_report
```

```
[119]:
```

	Not Fraud	Fraud
precision	0.596491	0.909722
recall	0.954386	0.415873
f1-score	0.734143	0.570806
support	285.000000	315.000000

```
[120]: show_roc_auc_score(ytrain, trainpreds, ytest, testpreds)
```

Train ROC-AUC score : 1.000 || Test ROC-AUC score : 0.685

```
[121]: show_cv_score(p_dt_model, xtrain, ytrain)
```

Train Cross validation Acc : 69.458 %

12 Pipelined Naive Bayes Model with Consensus

```
[122]: base_model = GaussianNB()
p_feature_names = [d_cols_final, t_cols_final]
p_model_names = ['d_predictor', 't_predictor']
p_nb_model = PipelinedModel(base_model, p_feature_names, p_model_names)

p_nb_model.fit(xtrain, ytrain.values.reshape(-1,))#.fit(xtrain, ytrain)

trainpreds = p_nb_model.predict(xtrain)
testpreds = p_nb_model.predict(xtest)

baseline_acc = namedtuple('Baseline_Accuracy', ['training', 'test'])
p_nb_baseline_acc = baseline_acc(training = p_nb_model.
    ↳get_baseline_accuracy(ytrain, trainpreds),
                                test = p_nb_model.get_baseline_accuracy(ytest,
    ↳testpreds))
show_accuracy(p_nb_baseline_acc)
```

Train Acc : 76.54 % || Test Acc : 74.67 %

```
[123]: p_nb_train_report = show_classification_report(ytrain, trainpreds)
p_nb_train_report
```

```
[123]:
```

	Not Fraud	Fraud
precision	0.688876	0.961424
recall	0.978601	0.546835
f1-score	0.808569	0.697149
support	1215.000000	1185.000000

```
[124]: p_nb_test_report = show_classification_report(ytest, testpreds)
p_nb_test_report
```

```
[124]:
```

	Not Fraud	Fraud
precision	0.655012	0.976608
recall	0.985965	0.530159
f1-score	0.787115	0.687243
support	285.000000	315.000000

```
[125]: show_roc_auc_score(ytrain, trainpreds, ytest, testpreds)
```

Train ROC-AUC score : 0.763 || Test ROC-AUC score : 0.758

```
[126]: show_cv_score(p_nb_model, xtrain, ytrain)
```

Train Cross validation Acc : 75.250 %

13 Pipelined SVM Model with Consensus

```
[127]: base_model = SVC(random_state=10)
p_feature_names = [d_cols_final, t_cols_final]
p_model_names = ['d_predictor', 't_predictor']
p_svm_model = PipelinedModel(base_model, p_feature_names, p_model_names)

p_svm_model.fit(xtrain, ytrain.values.reshape(-1,))#.fit(xtrain, ytrain)

trainpreds = p_svm_model.predict(xtrain)
testpreds = p_svm_model.predict(xtest)

baseline_acc = namedtuple('Baseline_Accuracy', ['training', 'test'])
p_svm_baseline_acc = baseline_acc(training = p_svm_model.
    ↳get_baseline_accuracy(ytrain, trainpreds),
                                test = p_svm_model.
    ↳get_baseline_accuracy(ytest, testpreds))
show_accuracy(p_svm_baseline_acc)
```

Train Acc : 55.79 % || Test Acc : 52.33 %

13.1 Probabilistic Analysis

```
[128]: p_svm_train_report = show_classification_report(ytrain, trainpreds)
p_svm_train_report
```

```
[128]:
```

	Not Fraud	Fraud
precision	0.534011	0.955882
recall	0.995062	0.109705
f1-score	0.695027	0.196821

```
support      1215.000000  1185.000000
```

```
[129]: p_svm_test_report = show_classification_report(ytest, testpreds)
p_svm_test_report
```

```
[129]:
```

	Not Fraud	Fraud
precision	0.499118	0.939394
recall	0.992982	0.098413
f1-score	0.664319	0.178161
support	285.000000	315.000000

```
[130]: show_roc_auc_score(ytrain, trainpreds, ytest, testpreds)
```

Train ROC-AUC score : 0.552 || Test ROC-AUC score : 0.546

```
[131]: show_cv_score(p_svm_model, xtrain, ytrain)
```

Train Cross validation Acc : 54.875 %

14 Pipelined Logistic Regression Model with Consensus

```
[132]: base_model = LogisticRegression(random_state=10)
p_feature_names = [d_cols_final, t_cols_final]
p_model_names = ['d_predictor', 't_predictor']
p_lr_model = PipelinedModel(base_model, p_feature_names, p_model_names)

p_lr_model.fit(xtrain, ytrain.values.reshape(-1,))

trainpreds = p_lr_model.predict(xtrain)
testpreds = p_lr_model.predict(xtest)

baseline_acc = namedtuple('Baseline_Accuracy', ['training', 'test'])
p_lr_baseline_acc = baseline_acc(training = p_lr_model.
    ↳get_baseline_accuracy(ytrain, trainpreds),
                                test = p_lr_model.
    ↳get_baseline_accuracy(ytest, testpreds))
show_accuracy(p_lr_baseline_acc)
```

Train Acc : 50.62 % || Test Acc : 47.50 %

14.1 Probabilistic Analysis

```
[133]: p_lr_train_report = show_classification_report(ytrain, trainpreds)
p_lr_train_report
```

```
[133]:
```

	Not Fraud	Fraud
precision	0.506250	0.0
recall	1.000000	0.0
f1-score	0.672199	0.0
support	1215.000000	1185.0

```
[134]: p_lr_test_report = show_classification_report(ytest, testpreds)
p_lr_test_report
```

```
[134]:
```

	Not Fraud	Fraud
precision	0.475000	0.0
recall	1.000000	0.0
f1-score	0.644068	0.0
support	285.000000	315.0

```
[135]: show_roc_auc_score(ytrain, trainpreds, ytest, testpreds)
```

Train ROC-AUC score : 0.500 || Test ROC-AUC score : 0.500

```
[136]: show_cv_score(p_lr_model, xtrain, ytrain)
```

Train Cross validation Acc : 50.625 %

```
[137]: p_metric_tuples = (
    (p_dt_baseline_acc.training, p_dt_baseline_acc.test),
    (p_nb_baseline_acc.training, p_nb_baseline_acc.test),
    (p_svm_baseline_acc.training, p_svm_baseline_acc.test),
    (p_lr_baseline_acc.training, p_lr_baseline_acc.test)
)

variant_2_results = compile_variant_results(p_metric_tuples, ['dt', 'nb',
    ↪ 'svm', 'lr'], only_baseline=True)
variant_2_results
```

```
[137]:
```

	baseline_train_acc	baseline_test_acc
dt	100.00	67.17
nb	76.54	74.67
svm	55.79	52.33
lr	50.62	47.50

```
[138]: #saving pipelined Model results
variant_2_results.to_csv(f'{REPORTS_DIR}/variant_2_baseline_performance.csv',
    ↪index = False)
```

15 Recommendation

The best model out of the 4 kinds built is the Naive Bayes model. It has a test accuracy of ~ 75 %.

16 Conclusion

Thus as per approach 1 (HCCFD Mixed Feature variant), Decision Tree is the best suited model while with approach 2 (HCCFD Pipelined Model with consensus driven inference), Naive Bayes proves to be the performant one.